

Pushti Shah

San Diego, CA | +1 (609) 349 0863 | pushah@ucsd.edu | linkedin.com/in/pushti11 | github.com/pushtishah11

EDUCATION

University of California, San Diego

*B.S. Mathematics and Computer Science, Minor in Data Science | GPA: 3.8
Tutor, DSC 198 (DSROCKS I)*

La Jolla, CA
March 2027
Fall 2026 – Present

WORK EXPERIENCE

Paramount (CBS News)

AI Engineering Intern

New York, NY
June 2026 – Aug 2026

- Took an ambiguous problem from archivists (which of 3.3M archived video assets are most at risk) to a production solution: audited the live database schema, turned their judgment into requirements, and designed and shipped a Python/PostgreSQL risk-scoring system the team still runs post-internship.
- Deployed an LLM extraction layer on AWS Bedrock pulling structured rights and compliance terms from free-text metadata that rules could not parse, and traced a precision failure to prompt design rather than model capability, letting a cheaper model outperform a costlier one.
- Cut full-archive scoring runtime 25x (110 minutes to under 5) by diagnosing a memory-bound allocation pattern and re-architecting the pipeline to stream metadata in constant-memory chunks; hardened it with a 48-test regression suite.
- Built an async video indexing pipeline on the Mimir platform calling TwelveLabs Pegasus across an 8PB archive, generating story breakdowns and IPTC news taxonomy classifications.

AutomaticGPT.ai

Software Engineer Intern

San Diego, CA
Sept 2025 – Dec 2025

- Built and shipped reusable React + Vite components for a chat-based AI workflow automation platform with 500+ active users, and integrated Auth0 and Stripe with backend REST APIs, including asynchronous payment failure paths.
- Fixed a customer-facing defect where conversation history with the AI assistant reset on tab switches and reloads by implementing chat persistence and client-side state management; shipped in Agile sprints with code reviews.

Youphoria (Wellness Startup)

Technical Product Manager

San Diego, CA
Sept 2024 – Sept 2025

- Built survey pipelines and SQL/Python data workflows to automate user data collection, cohort segmentation, and behavioral analysis, cutting manual research effort by ~40%.
- Developed personalization logic using recommendation heuristics and behavioral clustering to generate tailored wellness insights and tracking features, contributing to ~25% growth in early user engagement.

Tenup Software Solutions

Machine Learning Intern

Gujarat, India
July 2025 – Sept 2025

- Built a Python NLP pipeline normalizing 10K+ noisy, duplicate platform skills into canonical entries with embeddings, HDBSCAN, and K-Means, tuning preprocessing and parameters against evaluation results.
- Shipped an incremental auto-promotion workflow integrating new clusters into the platform taxonomy, cutting manual updates by 40%; documented the pipeline and presented evaluation reports to the team.

PROJECTS

RAG Search Engine for Mortgage Document Parsing | *Python, RAG, LlamaIndex, Faiss, OCR*

- Built a modular RAG pipeline over 200+ page mortgage documents: Tesseract and PaddleOCR extraction, PyMuPDF structured parsing, automated document classification, and Faiss + LlamaIndex retrieval.
- Cut document processing latency 15% while reaching zero errors on key document insights.

Networking Follow-Up Automation | *n8n, Google Gemini, Gmail API, Google Sheets API*

- Built an AI workflow that turns meeting notes into a personalized follow-up with Gemini and a structured output parser, logs the contact to Sheets, and saves a Gmail draft so a human reviews before anything is sent.

SKILLS

Languages: Python, SQL, Java, C++, C, JavaScript

Applied AI: RAG, LLM and Prompt Evaluation, AWS Bedrock, LlamaIndex, Faiss, Embeddings, Gemini, n8n, TwelveLabs, OCR

Cloud & Infrastructure: AWS (Bedrock, S3, Lambda, Secrets Manager), Docker, Git, CI/CD, Linux

Backend & Data: PostgreSQL, REST APIs, React, Express.js, Pandas, NumPy, Scikit-learn, Regression Testing